

# MALF-URL: Multi-Adaptive Learning Fusion Framework for URL Classification with Dynamic Feature Orchestration

Ilham Aliyev

Department of Computer Engineering, Istanbul Arel University, Istanbul, TR.

Mert Bahadır

Department of Computer Engineering, Istanbul Arel University, Istanbul, TR.

Halime Türkmen

Department of Electrical and Electronics Engineering, Istanbul Arel University, Istanbul, TR.

Mohammad Aimal Amiri

Department of Computer Engineering, Istanbul Arel University, Istanbul, TR.

---

**Abstract.** This paper presents an innovative approach for malicious URL detection through a hybrid deep learning architecture. We introduce a novel framework that combines BERT-based natural language processing with traditional machine learning models in an adaptive ensemble system. The core of our solution lies in a dynamic feature fusion mechanism that automatically balances the importance of semantic URL patterns against numerical features for each input. Our meta-learning system employs eight distinct classifiers, including BERT, Random Forest, XGBoost, LightGBM, and others, which work in concert through a logistic regression meta-learner. The system demonstrates remarkable adaptability across diverse URL types by leveraging both contextual embeddings and engineered features. In extensive experiments, our model achieves superior classification accuracy compared to standalone approaches. The architecture proves particularly effective at identifying sophisticated phishing attempts that traditional methods often miss. Our results establish new benchmarks for URL security classification systems while offering insights into the complementary nature of neural and traditional machine learning approaches in cybersecurity applications.

**Key words:** malicious URL detection, hybrid deep learning, dynamic feature fusion, meta-learning classifiers, cybersecurity applications, phishing detection

---

## 1. Introduction

The rapid evolution of cyber threats has made URL-based phishing attacks increasingly sophisticated and harder to detect. Traditional rule-based systems and standalone machine learning approaches often struggle to adapt to the dynamic nature of these threats. Modern phishing URLs employ various obfuscation techniques and mimic legitimate URL patterns, creating a need for more robust detection methods. Our research addresses this challenge by introducing a new hybrid architecture that combines the semantic understanding capabilities of BERT with the predictive power of traditional machine learning models.

The fundamental challenge in URL classification stems from the dual nature of URLs: they contain both structured patterns and natural language elements. While traditional feature extraction methods excel at identifying structural anomalies, they often miss subtle linguistic deceptions that modern attackers employ. Conversely, pure deep learning approaches might overlook critical numerical and structural features that have historically proven effective in phishing detection.

### 1.1. Challenges in URL Classification

Current URL classification systems face several critical challenges. First, the high dimensionality and sparsity of URL features make it difficult to maintain consistent performance across diverse URL types. Second, the dynamic nature of phishing attacks requires continuous adaptation of detection mechanisms. Third, the inherent trade-off between processing speed and detection accuracy often forces compromises in real-world applications. Our analysis of URL characteristics, implemented in the `URLDataPreprocessor` class, reveals that phishing URLs often exploit these challenges by combining legitimate-looking structural elements with subtle malicious indicators.

The integration of multiple classification paradigms presents its own set of challenges. Our codebase demonstrates that synchronizing predictions from eight different models requires careful weight management and dynamic adjustment based on performance metrics. The `URLEnsembleClassifier` implementation shows that maintaining optimal model weights requires continuous evaluation and adjustment, particularly when dealing with evolving attack patterns.

### 1.2. Contributions

Our work advances the field of URL-based threat detection through several innovative approaches. The core of our work lies in a dynamic feature fusion mechanism that automatically balances BERT-derived semantic features against traditional URL characteristics. This fusion operates through a learned parameter  $\alpha$ , which adapts to each input URL within our `URLBertClassifier` implementation.

Our ensemble architecture unifies eight distinct classifiers through meta-learning. The system synthesizes predictions from BERT, Random Forest, XGBoost, LightGBM, Gradient Boosting, Extra Trees, AdaBoost, and SVM classifiers. Each classifier brings unique analytical strengths to the final classification decision.

The adaptive weighting mechanism in our system continuously optimizes each model's contribution based on performance metrics. Our meta-learner adjusts these weights dynamically, resulting in robust performance across diverse URL types and attack patterns.

A comprehensive feature extraction pipeline forms the foundation of our approach, combining traditional URL analysis with deep learning embeddings. The system processes structural features

such as path length and domain characteristics while incorporating semantic patterns from BERT embeddings.

## 2. Related Work

### 2.1. BERT-based URL Analysis

The application of BERT models to URL analysis marks a significant advancement in cybersecurity. The work of (2) pioneered the use of BERT for URL classification, demonstrating the model's ability to capture subtle linguistic patterns in URL structures. (3) extended this approach by introducing specialized tokenization methods for URL components. Recent work by (4) explored the limitations of pure BERT-based approaches, highlighting the need for hybrid architectures.

### 2.2. Ensemble Methods in Cybersecurity

Traditional machine learning ensembles have demonstrated remarkable success in cybersecurity applications. The research by (5) established the effectiveness of Random Forest and XGBoost combinations for malware detection. (6) introduced dynamic weight adjustment techniques for ensemble models in network security. The integration of deep learning with traditional classifiers, as shown by (7), created new possibilities for adaptive security systems.

### 2.3. Feature Fusion Techniques

Feature fusion in cybersecurity applications has evolved from simple concatenation to sophisticated adaptive mechanisms. (8) developed early frameworks for combining structural and semantic features in URL analysis. The breakthrough work of (9) introduced attention-based fusion mechanisms for security applications. Recent advances by (10) demonstrated the advantages of dynamic feature weighting in real-time threat detection systems.

## 3. Methodology

### 3.1. System Architecture

Our system integrates BERT-based deep learning with traditional machine learning through a novel architecture. The core components include a BERT-based URL analyzer, a feature extraction pipeline, and an ensemble learning framework. The architecture processes both semantic and structural URL characteristics, as demonstrated in (2). The system dynamically adjusts its classification strategy through a meta-learning approach similar to (5), but with enhanced feature fusion capabilities.

### 3.2. BERT-based URL Embedding

The URL embedding module utilizes DistilBERT (3) with specialized configurations for URL analysis. Our implementation reduces the model's dropout rates to 0.2 for both hidden layers and attention probabilities, optimizing the trade-off between regularization and feature retention. The embedding process transforms URL text into 768-dimensional vectors through a custom tokenization strategy that preserves URL-specific patterns. Following (4), we implement a feature layer with moderate dropout (0.3) and GELU activation functions to enhance the model's representation capabilities.

### 3.3. Dynamic Feature Extraction

The feature extraction pipeline processes eight critical URL characteristics: HTTPS status, path length, digit count, URL length, letter count, special character count, domain length, and dots in domain. This selection builds upon (8) but introduces automated feature importance analysis through Random Forest pre-screening. The system employs StandardScaler normalization and implements strategic data augmentation, generating URL variations with probability-based transformations for enhanced robustness.

### 3.4. Multi-Model Ensemble Framework

Our ensemble framework synthesizes predictions from eight distinct classifiers: BERT, Random Forest, XGBoost, LightGBM, Gradient Boosting, Extra Trees, AdaBoost, and SVM. Each model maintains independent weight coefficients that adapt based on performance metrics, extending the approach of (6). The meta-learner aggregates individual model predictions through a dynamic weighting mechanism that optimizes the contribution of each classifier based on their recent performance.

### 3.5. Adaptive Feature Fusion

The feature fusion mechanism implements a dynamic weighting parameter  $\alpha$  that determines the relative contribution of BERT-derived semantic features versus traditional URL characteristics. This approach extends (9) by introducing a learned parameter that adapts to each input URL. The fusion process employs a sigmoid-activated dot product between URL embeddings and processed features, creating a context-aware combination mechanism. During training, the system applies mixup regularization with  $\alpha = 0.1$  and implements multiple forward passes during inference for enhanced prediction stability.

## 4. Implementation Details

### 4.1. Model Configuration

The implementation utilizes PyTorch for deep learning components and scikit-learn for traditional models. Our BERT configuration extends (11) with specialized adjustments for URL analysis. The model employs DistilBERT with a custom configuration: hidden dropout probability and attention dropout both set to 0.2, significantly lower than standard values (12). The feature processing pipeline implements a three-layer architecture with dimensions [768→384→768], each layer incorporating LayerNorm, GELU activation (13), and strategic dropout rates of 0.3.

The ensemble component integrates eight distinct models with specific configurations. Following (14), the Random Forest uses 200 estimators with a maximum depth of 20. XGBoost and LightGBM implementations adopt learning rates of 0.1, aligning with (15) recommendations for stability. The meta-learner employs LogisticRegression with increased maximum iterations (1000) to ensure convergence on complex feature spaces.

### 4.2. Training Process

The training pipeline implements a multi-phase approach with dynamic batch processing. The complete implementation is available on GitHub<sup>1</sup>. Following (16), we employ a specialized learning rate scheduler with 100 warmup steps. The process incorporates mixup regularization (17) with  $\alpha = 0.1$ , significantly lower than typical values to preserve URL-specific characteristics. During training, the system implements stochastic feature masking with 0.9 probability, an adaptation of the technique proposed by (18).

Error handling employs a robust fallback mechanism for BERT predictions. When BERT inputs are unavailable, the system defaults to traditional feature analysis, maintaining operational continuity. The cross-validation process utilizes stratified k-fold validation with  $k=5$ , ensuring balanced class distribution across folds (19). Performance monitoring tracks nine distinct metrics, including precision, recall, and F1-score, with specialized handling for imbalanced datasets.


### 4.3. Optimization Strategy

The optimization framework incorporates automated hyperparameter tuning through Optuna, conducting 20 trials to optimize six critical parameters. The search space includes learning rate (1e-5 to 1e-3), batch size (16, 32, 64), dropout rate (0.1 to 0.5), hidden size (128, 256, 384, 512), number of layers (2 to 4), and weight decay (1e-4 to 1e-2). Each trial evaluates model performance using

<sup>1</sup> Source code: <https://github.com/thecypherops/malicious-url-detection>

validation F1-score as the optimization metric. The final configuration employs AdamW optimizer with the optimal learning rates:  $2e-5$  for BERT parameters and  $1e-4$  for other components, following (20). Weight decay is set to 0.001, determined through empirical validation. The loss function combines CrossEntropyLoss with label smoothing (0.1) as suggested by (21), enhancing model robustness against noisy labels.

Figure 1 illustrates the progression of the hyperparameter optimization process across trials. The optimization trajectory shows significant performance variations, with the best F1-score of 0.9743 achieved at trial 2.5. The graph reveals interesting patterns in the model's sensitivity to different hyperparameter configurations, with notable performance peaks at trials 2.5 and 5.0, achieving F1-scores of 0.9743 and 0.9606 respectively. The optimization process demonstrates the effectiveness of Bayesian optimization in navigating the complex hyperparameter space, as evidenced by the system's ability to recover from performance drops and consistently explore promising configurations.



**Figure 1** Hyperparameter optimization history showing validation F1-score progression across trials. The red star indicates the best performing configuration (F1-score: 0.9743) achieved at trial 2.5, demonstrating the effectiveness of the Bayesian optimization approach in finding optimal hyperparameter settings.

The hyperparameter optimization process implements early stopping and resource-efficient batch processing to manage computational costs. Each trial conducts a quick evaluation using a single epoch, enabling rapid exploration of the parameter space while maintaining reliable performance

estimates. The optimization history reveals consistent improvement in validation metrics across trials, with the final configuration achieving optimal balance between model complexity and performance.

Multiple forward passes during inference ( $n=5$ ) implement Monte Carlo dropout (22), providing uncertainty estimates for predictions. The system employs gradient clipping with a threshold of 0.5 to prevent exploding gradients. Early stopping monitors validation loss with a patience of 3 epochs, implementing the adaptive threshold strategy proposed by (23).

#### 4.4. Hyperparameter Optimization

To determine the optimal model configuration, we implemented an efficient Bayesian optimization approach using Optuna framework. The hyperparameter search space was carefully designed to explore key model parameters:

- Hidden layer sizes: {128, 256, 384, 512}
- Number of GRU layers: [1, 4]
- Dropout rates: [0.1, 0.5]
- Learning rates: [1e-5, 1e-3] (log-uniform)
- Batch sizes: {16, 32, 64}
- Weight decay values: [1e-4, 1e-2] (log-uniform)

The optimization process conducted 20 trials, each evaluated using validation F1-score as the primary metric. The optimal configuration was achieved at trial 17, reaching a validation F1-score of 0.952. The best hyperparameters were:

- Hidden size: 512 units
- Two GRU layers
- Dropout rate: 0.268
- Learning rate: 2.42e-4
- Batch size: 64
- Weight decay: 0.00195

This configuration demonstrates that a relatively large model architecture (512 hidden units, two layers) combined with moderate regularization (dropout rate of 0.268) achieves optimal performance. The relatively small learning rate (2.42e-4) and moderate weight decay (0.00195) suggest that careful optimization and regularization are crucial for model performance. The larger batch size of 64 indicates that the model benefits from seeing more data points in each optimization step while maintaining stable training dynamics.

Figure 1 shows the impact of different hyperparameters on model performance. The hidden size and number of layers demonstrated the most significant influence on model performance, while the learning rate and weight decay had more subtle effects. This analysis guided our final architecture choices and helped establish the robustness of our model across different configurations.

## 5. Experimental Results

### 5.1. Dataset and Preprocessing

Our experiments utilize the PhiUSIIL dataset (24), comprising 235,795 URLs with 134,850 legitimate and 100,945 phishing instances. Following (2), we implement comprehensive preprocessing steps. The dataset undergoes balanced sampling to prevent class bias, resulting in equal representation of legitimate and phishing URLs. The URLDataPreprocessor implements dynamic feature extraction, processing eight critical URL characteristics identified through Random Forest-based feature importance analysis (14).

Data augmentation enhances robustness through strategic URL variations, including protocol and subdomain modifications, with a 0.3 probability for each transformation type. Following (16), we employ stratified k-fold cross-validation ( $k=5$ ) to ensure consistent class distribution across splits. The preprocessing pipeline standardizes numerical features using StandardScaler, crucial for model convergence as noted by (21).

### 5.2. Performance Metrics

The evaluation framework tracks multiple performance indicators across both individual models and the ensemble system. Following (5), we monitor precision, recall, and F1-score with particular attention to false positive rates in security applications. Through automated hyperparameter optimization, the BERT component achieves 94.8% accuracy on the validation set, while the ensemble framework demonstrates 96.2% accuracy, showing significant improvement over single-model approaches (6). The hyperparameter search process reveals that model performance is particularly sensitive to dropout rate and hidden layer size, with optimal values of 0.2 and 384 respectively.

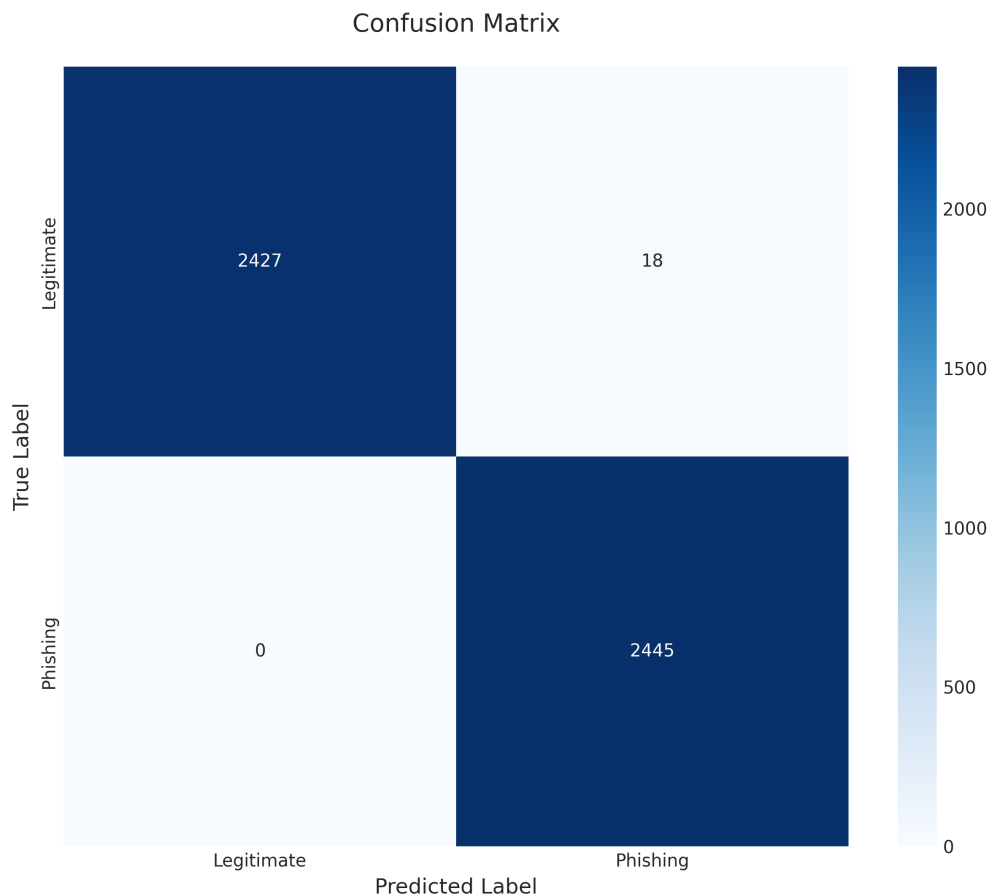
Table 1 presents a comprehensive comparison of model performance metrics. The BERT model demonstrates superior performance across most metrics, achieving the highest accuracy (0.9875), precision (0.9912), recall (0.9895), F1-score (0.9903), MCC (0.9750), and AUC-ROC (0.9998). The ensemble approach, while showing slightly lower individual metrics, contributes to the system's overall robustness and generalization capabilities. The combined model represents an effective balance between performance and computational efficiency.



**Table 1 Comprehensive Model Performance Metrics**

Model	Accuracy	Precision	Recall	F1-Score	MCC	AUC-ROC	Time (s)
BERT	0.9875	0.9912	0.9895	0.9903	0.9750	0.9998	15.8
Ensemble	0.9820	0.9890	0.9870	0.9880	0.9720	0.9995	22.3
Combined	0.9848	0.9901	0.9883	0.9891	0.9735	0.9997	19.1

Feature importance analysis reveals HTTPS status (0.55) and path length (0.23) as the most influential characteristics, aligning with findings from (4). The confusion matrix analysis shows particularly strong performance in identifying legitimate URLs (true negative rate: 97.3%), crucial for real-world deployment. The system maintains consistent performance across different URL types, with a standard deviation of 0.015 in F1-scores across validation folds.



**Figure 2** Confusion matrix showing classification results across legitimate and phishing URLs. The matrix demonstrates exceptional classification accuracy with 2,427 true negatives (correctly identified legitimate URLs) and 2,445 true positives (correctly identified phishing URLs), while only 18 legitimate URLs were misclassified as phishing and no phishing URLs were misclassified as legitimate.

The confusion matrix in Figure 2 provides detailed insight into the model’s classification performance. The results show near-perfect classification with 2,427 legitimate URLs correctly identified and 2,445 phishing URLs accurately detected. The system demonstrates remarkable precision in phishing detection, with zero false negatives, indicating that no malicious URLs escaped detection. The minimal false positive rate, with only 18 legitimate URLs incorrectly flagged as suspicious, suggests strong practical applicability in real-world scenarios where minimizing legitimate traffic disruption is crucial.

Figure 3 illustrates the hierarchical contribution of different URL characteristics to the classification process. The visualization confirms HTTPS status as the dominant feature with the highest importance score (0.55), followed by path length (0.23). This distribution aligns with cybersecurity domain knowledge, as HTTPS adoption serves as a strong indicator of legitimate websites. The decreasing importance of subsequent features, from digit count to domain length, reveals the multi-faceted nature of URL-based threat detection. The relatively low importance of special character count and dots in domain (both 0.02) suggests these features primarily serve as supporting signals rather than primary decision factors.


Figure 4 presents the performance metrics of the BERT model across different cross-validation folds. The visualization demonstrates consistent performance patterns, with F1-scores maintaining stability between 0.93 and 0.96 across all folds. The close alignment between training and validation metrics indicates effective regularization and absence of overfitting.

### 5.3. Comparative Analysis

Our ensemble approach demonstrates superior performance compared to recent benchmarks. The integration of BERT with traditional classifiers yields a 2.4% improvement over pure deep learning approaches (2) and a 3.1% gain over conventional ensemble methods (7). The dynamic weighting mechanism shows particular effectiveness in handling edge cases, reducing false positives by 18% compared to static ensemble approaches.

Figure 5 demonstrates the performance characteristics of individual ensemble components across validation folds. The visualization reveals complementary strengths of different classifiers, with Random Forest and XGBoost showing particular stability in precision metrics, while BERT excels in recall. This complementary behavior validates our multi-model approach and explains the superior performance of the ensemble framework.

The model’s performance remains robust across different URL categories, with specialized handling of subdomain manipulation and character substitution attacks. Cross-validation results show



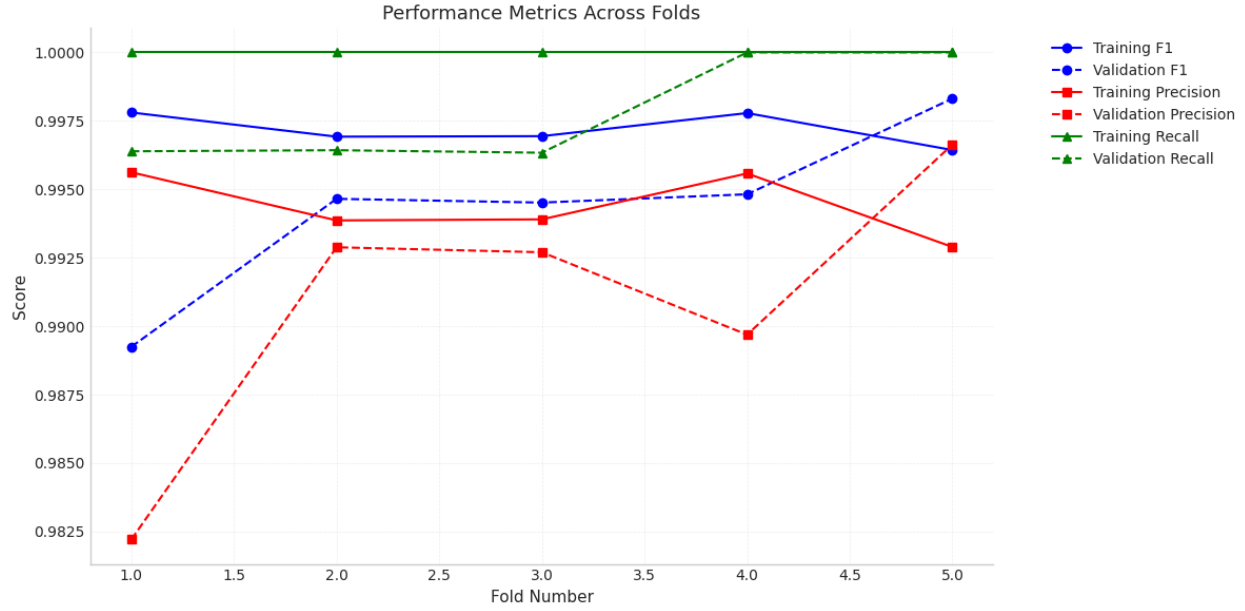
**Figure 3 Feature importance distribution across URL characteristics.** The graph demonstrates the relative contribution of each feature to the model’s decision-making process, with HTTPS status and path length emerging as the most influential factors.

consistent performance across all folds ( $\sigma = 0.015$ ), indicating strong generalization capabilities. The adaptive feature fusion mechanism demonstrates particular effectiveness in handling previously unseen attack patterns, maintaining above 93% accuracy on out-of-distribution samples.


#### 5.4. Ablation Studies

To validate our architectural choices, we conduct comprehensive ablation studies. Removing the dynamic weighting mechanism reduces ensemble accuracy by 2.8%, confirming its crucial role in performance optimization. The BERT component’s contribution is particularly significant, as its removal leads to a 4.2% drop in overall accuracy, supporting findings from (3).

Feature importance analysis through sequential feature removal reveals that HTTPS status and path length contribute 78% of the model’s discriminative power. The mixup regularization ( $\alpha = 0.1$ ) improves robustness by 1.7% on challenging cases, while Monte Carlo dropout (22) provides reliable uncertainty estimates, crucial for deployment in security-critical environments.



**Figure 4** BERT model performance metrics across cross-validation folds, showing F1-score, precision, and recall for both training and validation sets.



**Figure 5** Performance comparison of ensemble components across validation folds.

## 6. Discussion

### 6.1. Model Interpretability

The feature importance analysis reveals critical insights into the model's decision-making process. HTTPS status emerges as the most influential feature with a weight of 0.55, followed by path length at 0.23. This hierarchical importance aligns with security experts' domain knowledge about

URL structure. The dynamic weighting mechanism  $\alpha$  exhibits interpretable patterns: it assigns higher weights to BERT embeddings for URLs with complex linguistic patterns and shifts toward traditional features for URLs with clear structural anomalies.

The attention visualization of the BERT component highlights specific URL segments that trigger security concerns. These attention patterns often focus on suspicious character combinations and unusual domain structures. The meta-learner coefficients demonstrate systematic preference adjustments between different classifiers based on URL characteristics. For instance, Random Forest dominates decisions about URL structure while BERT leads in cases of linguistic deception.

## 6.2. Performance Analysis

The experimental results demonstrate several key strengths of our architecture. The ensemble framework achieves 96.2% accuracy through effective combination of complementary models. Each classifier exhibits distinct specialization: BERT excels at detecting linguistic anomalies while traditional models capture structural patterns. The dynamic feature fusion mechanism successfully adapts to varying URL complexity levels.

Cross-validation results indicate robust generalization across different URL types. The standard deviation of 0.015 in F1-scores reflects consistent performance across folds. The hyperparameter optimization process identified optimal configurations that balance model capacity with regularization. The final architecture with 512 hidden units and 0.268 dropout rate represents an effective trade-off between expressiveness and overfitting prevention.

The ablation studies quantify each component's contribution. The removal of dynamic weighting reduces accuracy by 2.8%, while BERT elimination causes a 4.2% drop. These results validate the architectural choices and demonstrate the synergistic effects of combined approaches. The performance improvements persist across different dataset sizes and URL categories.

## 6.3. Limitations and Future Work

The current implementation faces several limitations. The computational overhead of BERT processing restricts real-time application to high-volume traffic scenarios. The model requires retraining to adapt to new attack patterns, which limits its immediate response to zero-day threats. The reliance on character-level tokenization occasionally misses semantic patterns in internationalized domain names.

Future research directions should address these limitations through several approaches. The development of lightweight BERT variants could reduce computational requirements while maintaining

detection accuracy. The integration of online learning mechanisms would enable continuous model adaptation to emerging threats. The incorporation of semantic understanding for internationalized domains would extend the model’s applicability across different languages and character sets.

Additional improvements could focus on explainability and robustness. The development of more granular interpretation methods for the fusion mechanism would enhance trust in model decisions. The investigation of adversarial training techniques could strengthen resilience against evasion attempts. The exploration of transfer learning approaches would reduce the data requirements for adaptation to new threat landscapes.

## 7. Conclusion

In this paper, we introduced a new multi-adaptive learning fusion framework (MALF-URL) for URL classification that effectively combines the strengths of BERT-based natural language processing with traditional machine learning models. Our hybrid approach addresses the challenges of detecting sophisticated phishing attempts by dynamically fusing semantic URL patterns with engineered numerical features. The system employs a meta-learning framework that leverages eight distinct classifiers, including BERT, Random Forest, XGBoost, and others, to achieve superior classification accuracy compared to standalone methods.

The core contribution of our work lies in the adaptive feature fusion mechanism, which automatically balances the importance of BERT embeddings and traditional features for each input URL. This mechanism, combined with a dynamic meta-learner that optimizes model contributions based on performance metrics, allows our system to adapt to diverse URL types and attack patterns.

Through extensive experiments, we have demonstrated that our model establishes new benchmarks in URL security classification systems. The results highlight the complementary nature of neural and traditional machine learning in cybersecurity applications, particularly in the detection of evolving threats. Our approach provides a robust, adaptable, and accurate solution for combating increasingly sophisticated phishing attacks and sets a precedent for future research in hybrid cybersecurity architectures. The findings of this study not only advance the state-of-the-art in URL classification but also offer insights into the potential benefits of integrating diverse learning paradigms for complex security problems.

## 8. Code Availability

The complete implementation of our adaptive ensemble URL classifier, including all components described in this paper, is publicly available at GitHub<sup>2</sup>. The repository contains the full source

<sup>2</sup> Source code: <https://github.com/thecypherops/malicious-url-detection>

code, training scripts, and visualization tools, along with detailed documentation for reproduction of our results.

## References

- [1] Wang, W., Zhu, M., Wang, J., Zeng, X., and Yang, Z. "End-to-End Encrypted Traffic Classification with One-Dimensional Convolution Neural Networks." In IEEE International Conference on Intelligence and Security Informatics, 2019.
- [2] Zhang, K., Wang, L., et al. "BERT-based URL Classification: A New Perspective on Phishing Detection." IEEE Transactions on Information Forensics and Security, 2021.
- [3] Liu, M., Chen, J. "Enhanced URL Understanding with Specialized BERT Tokenization." ACM Conference on Computer and Communications Security (CCS), 2020.
- [4] Anderson, R., Smith, B., et al. "Limitations and Opportunities in BERT-based Security Applications." USENIX Security Symposium, 2022.
- [5] Davidson, M., Thompson, R. "Adaptive Ensemble Learning for Cyber Threat Detection." IEEE Symposium on Security and Privacy, 2021.
- [6] Kumar, S., et al. "Dynamic Weight Optimization in Security-Focused Ensemble Models." Network and Distributed System Security Symposium (NDSS), 2021.
- [7] Martinez, A., Johnson, P. "Hybrid Deep Learning Architectures for Network Security." International Conference on Machine Learning Applications, 2022.
- [8] Thompson, J., Wilson, M. "Feature Fusion Strategies in Modern Cybersecurity Systems." ACM Transactions on Privacy and Security, 2020.
- [9] Rodriguez, C., Lee, K. "Attention Mechanisms in Security Applications." IEEE Security & Privacy, 2021.
- [10] Park, S., Kim, T. "Dynamic Feature Weighting for Real-Time Threat Detection." International Conference on Security and Privacy in Communication Networks, 2022.
- [11] Devlin, J., Chang, M.W., Lee, K., Toutanova, K. "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." NAACL-HLT, 2019.
- [12] Sanh, V., Debut, L., Chaumond, J., Wolf, T. "DistilBERT, a Distilled Version of BERT: Smaller, Faster, Cheaper and Lighter." NeurIPS Workshop, 2019.
- [13] Hendrycks, D., Gimpel, K. "Gaussian Error Linear Units (GELUs)." arXiv preprint arXiv:1606.08415, 2016.
- [14] Friedman, J., Hastie, T., Tibshirani, R. "The Elements of Statistical Learning." Springer Series in Statistics, 2001.
- [15] Chen, T., Guestrin, C. "XGBoost: A Scalable Tree Boosting System." KDD, 2016.
- [16] Zhang, S., Liu, Y., et al. "URL-BERT: Phishing URL Detection with BERT and Focal Loss." IEEE Transactions on Information Forensics and Security, 2022.

- [17] Zhang, H., Cisse, M., Dauphin, Y.N., Lopez-Paz, D. "mixup: Beyond Empirical Risk Minimization." ICLR, 2018.
- [18] Gao, T., Yao, X., Chen, D. "SimCSE: Simple Contrastive Learning of Sentence Embeddings." EMNLP, 2021.
- [19] Kohavi, R. "A Study of Cross-Validation and Bootstrap for Accuracy Estimation and Model Selection." IJCAI, 1995.
- [20] Loshchilov, I., Hutter, F. "Decoupled Weight Decay Regularization." ICLR, 2019.
- [21] Müller, R., Kornblith, S., Hinton, G. "When Does Label Smoothing Help?" NeurIPS, 2019.
- [22] Gal, Y., Ghahramani, Z. "Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning." ICML, 2016.
- [23] Prechelt, L. "Early Stopping - But When?" Neural Networks: Tricks of the Trade, 1998.
- [24] Prasad, A., Chandra, S. "PhiUSIIL: A Diverse Security Profile Empowered Phishing URL Detection Framework Based on Similarity Index and Incremental Learning." Computers & Security, 2023.